UNIVERSITY *of* WASHINGTON

LIMEaid
Local Interpretable Model-agnostic Explanations (LIME)

Data 515, Spring 2019
M.S. Data Science

Suman Bhagavathula
Patrick King
Javier Salido

Interpretability in Machine Learning

Some highly accurate models are not "explainable"

- Neural networks, random forests

Why is this a problem?

- Bias, not obvious

- High test set accuracy but poor results in the field

- Policy or law demands an explanation of any decision

Solution: model-agnostic local explanations

- Explain one instance, not entire model

- Fit a simple model to explain a small section of decision space

LIMEaid: A LIME solution for tabular data

LIMEaid explanations

Input

- A "complex" ML model, fit by sklearn classifier object with .predict
- An instance of data (x) and its model output ($f(x)$)
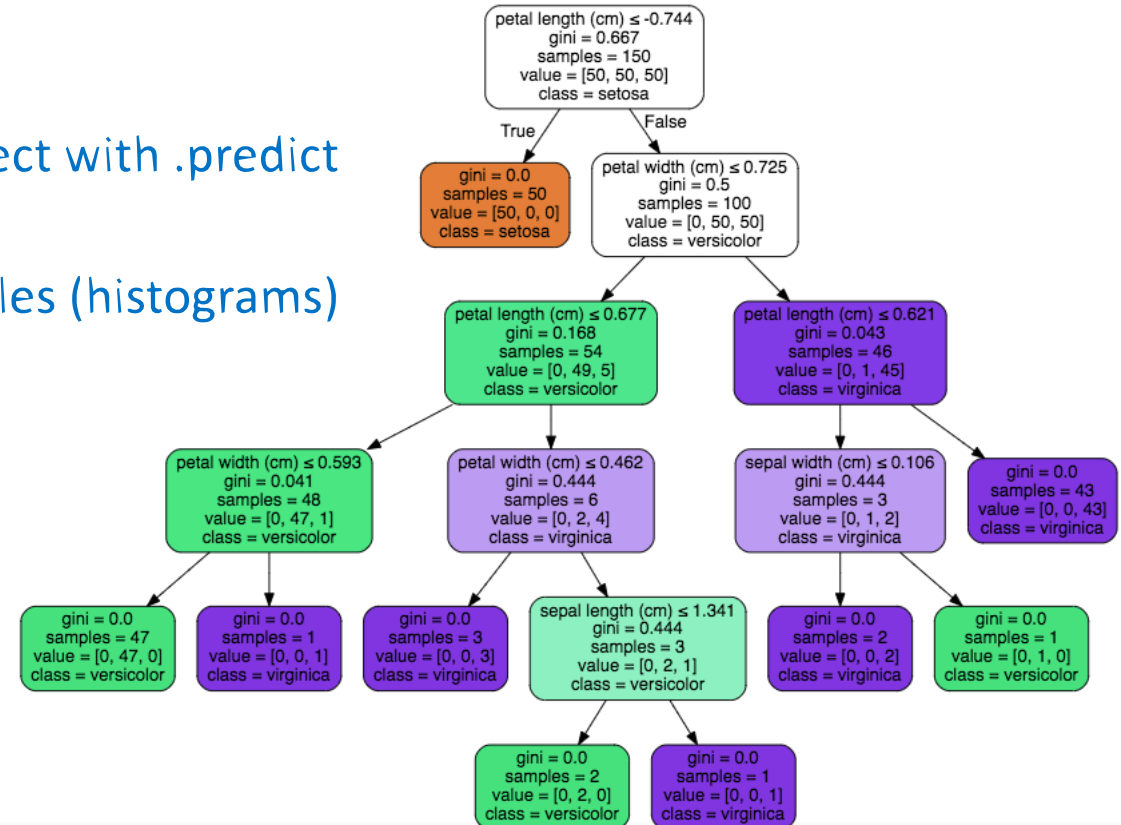- Probability domain for normalized predictor variables (histograms)

Output

- Sparse linear models (few features), plottable
  List significant features

Analysis/verification

- Decision-tree comparison

Use Cases

**User profile**: data scientist with Python programming experience

1. **Model verification scenario**
   - User wants to preempt poor model performance "in the field"
   - Use LIMEaid to sample test dataset
   - Show most significant features for decision
   - Tune or replace model if spurious correlation or other issues

2. **Decision explanation scenario**
   - Classification has already been made by a model
   - Use LIMEaid to sample whole dataset
   - Produce easy-to-share "two-dimensional" plot of a linear correlation

Design
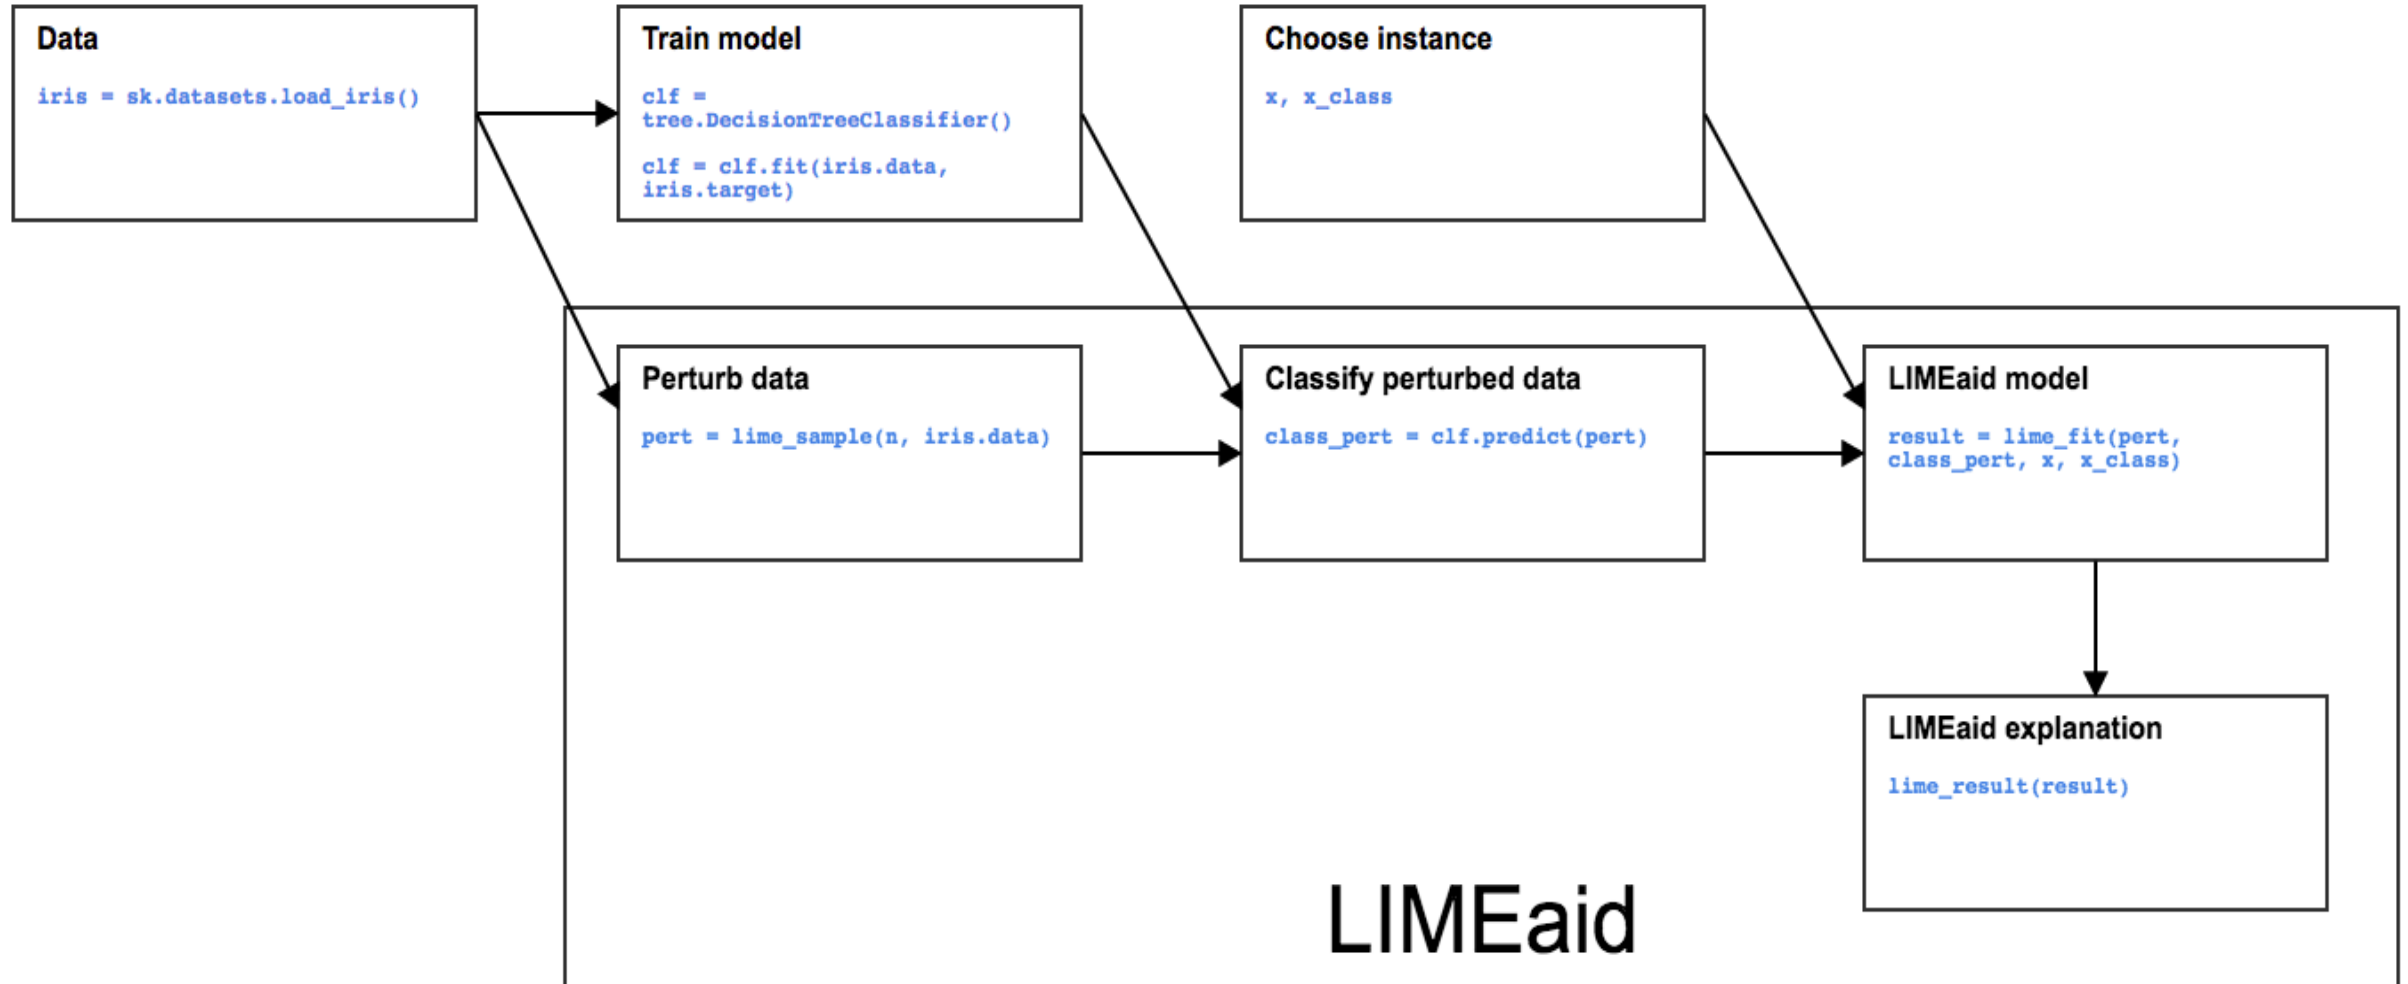
Data

Sources

- College Scorecard (data.gov): Annual report of schools and attributes (SAT scores, majors offered, region, cost, public/private/for profit, etc.)

- "Where it Pays to Attend College" (Kaggle.com) obtained from (*Wall Street Journal*), based on Payscale, Inc. (College Salary Report Methodology): Article reporting schools and salaries of graduates, salaries by major, etc.

Merge

- Significant cleaning, reformatting to match sets on college name

- String manipulation, removal of hyphens, abbreviations, region names, etc.

More

- Also tested with Sklearn's provided "Iris" data, to show comparison to Ribeiro's original LIME package ("LIME classic")

Models

Scikit-learn classifiers that predict probabilities (predict_proba implemented)

Multiclass logistic regression (sklearn.linear_model.LogisticRegression)
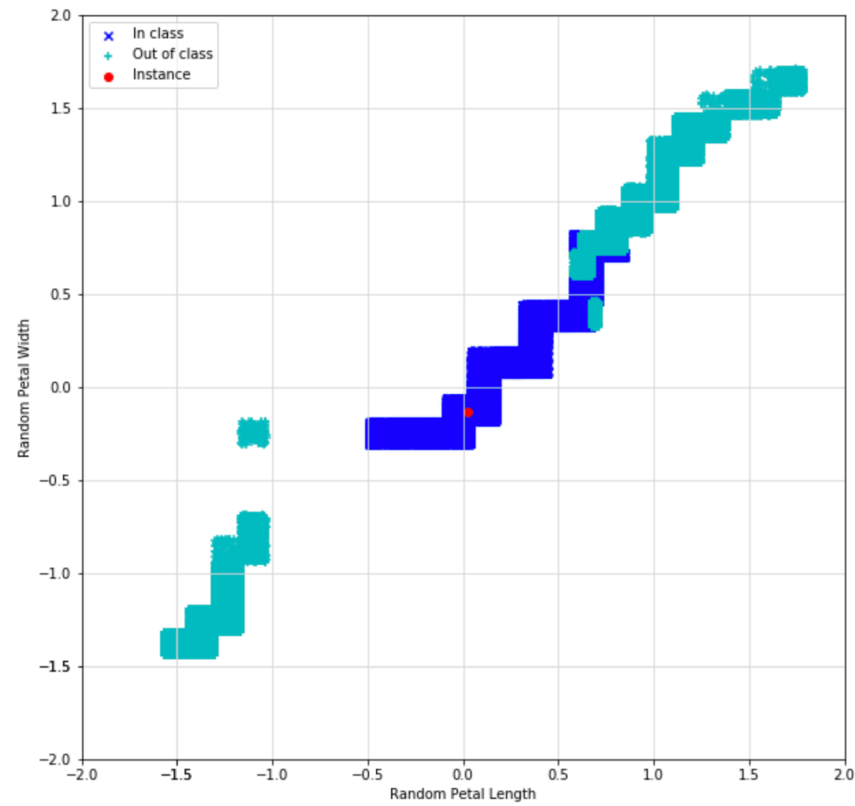- 85% accuracy on College data

Random Forest (sklearn.ensemble.randomforestclassifier)
- 65% accuracy on College data

- Decision tree (sklearn.tree)
  65% accuracy on College data

Models not tuned for improved accuracy (default settings)

Demo

Iris notebook (LIME_Iris_ex_notebook.ipynb)
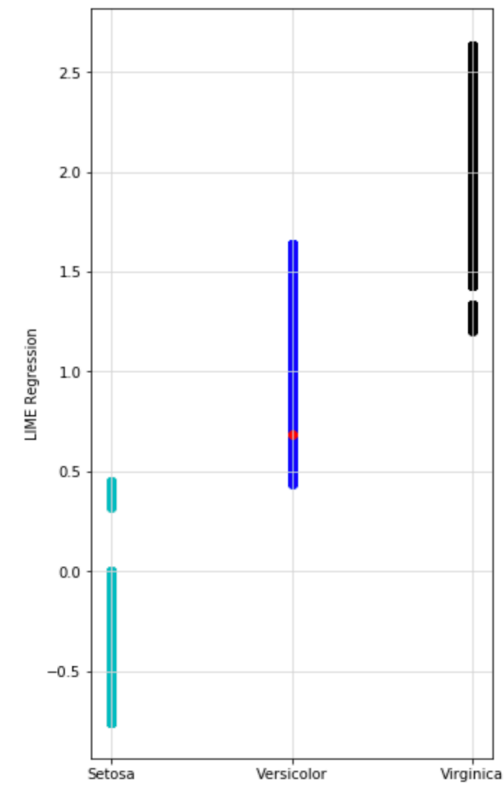
Project Structure

- [LIMEaid Github Repo](#) based on [Shablona](#)
- Used "M-V-C" architecture to codebase (Model-View-Controller)

Future Work

- **API support** for data acquisition to support dynamic features:
  - College Scorecard (data.gov) currently published with new features and data dictionary yearly

- **Model tuning** for examples:
  - Currently using defaults, could improve accuracy > 65%

- **Modify penalty** for number of coefficients

- **More data type and model support:** image data, NLP support, support for model objects beyond sklearn classifiers

Questions?