UNIVERSITY of WASHINGTON

LIMEaid

Local Interpretable Model-agnostic Explanations (LIME)

Data 515, Spring 2019
M.S. Data Science

Suman Bhagavathula
Patrick King
Javier Salido

Interpretability in Machine Learning

Some highly accurate models are not “explainable”

- Neural networks, random forests

Why is this a problem?

- Bias, not obvious
- High test set accuracy but poor results in the field
- Policy or law demands an explanation of any decision

Solution: model-agnostic local explanations

- Explain one instance, not entire model
- Fit a simple model to explain a small section of decision space

LIMEaid: A LIME solution for tabular data

Input

- A “complex” ML model, fit by sklearn classifier object with .predict
- An instance of data (*x*) and its model output (*f(x)*)
- Probability domain for normalized predictor variables (histograms)

Output

- Sparse linear models (few features), plottable
- List significant features

Install

- PyPI package

Use Cases

User profile: data scientist with Python programming experience

1. **Model verification scenario**
   - User wants to preempt poor model performance “in the field”
   - Use LIMEaid to sample test dataset
   - Show most significant features for decision
   - Tune or replace model if spurious correlation or other issues

2. **Decision explanation scenario**
   - Classification has already been made by a model
   - Use LIMEaid to sample whole dataset
   - Produce easy-to-explain “two-dimensional” plot of linear regression

Design

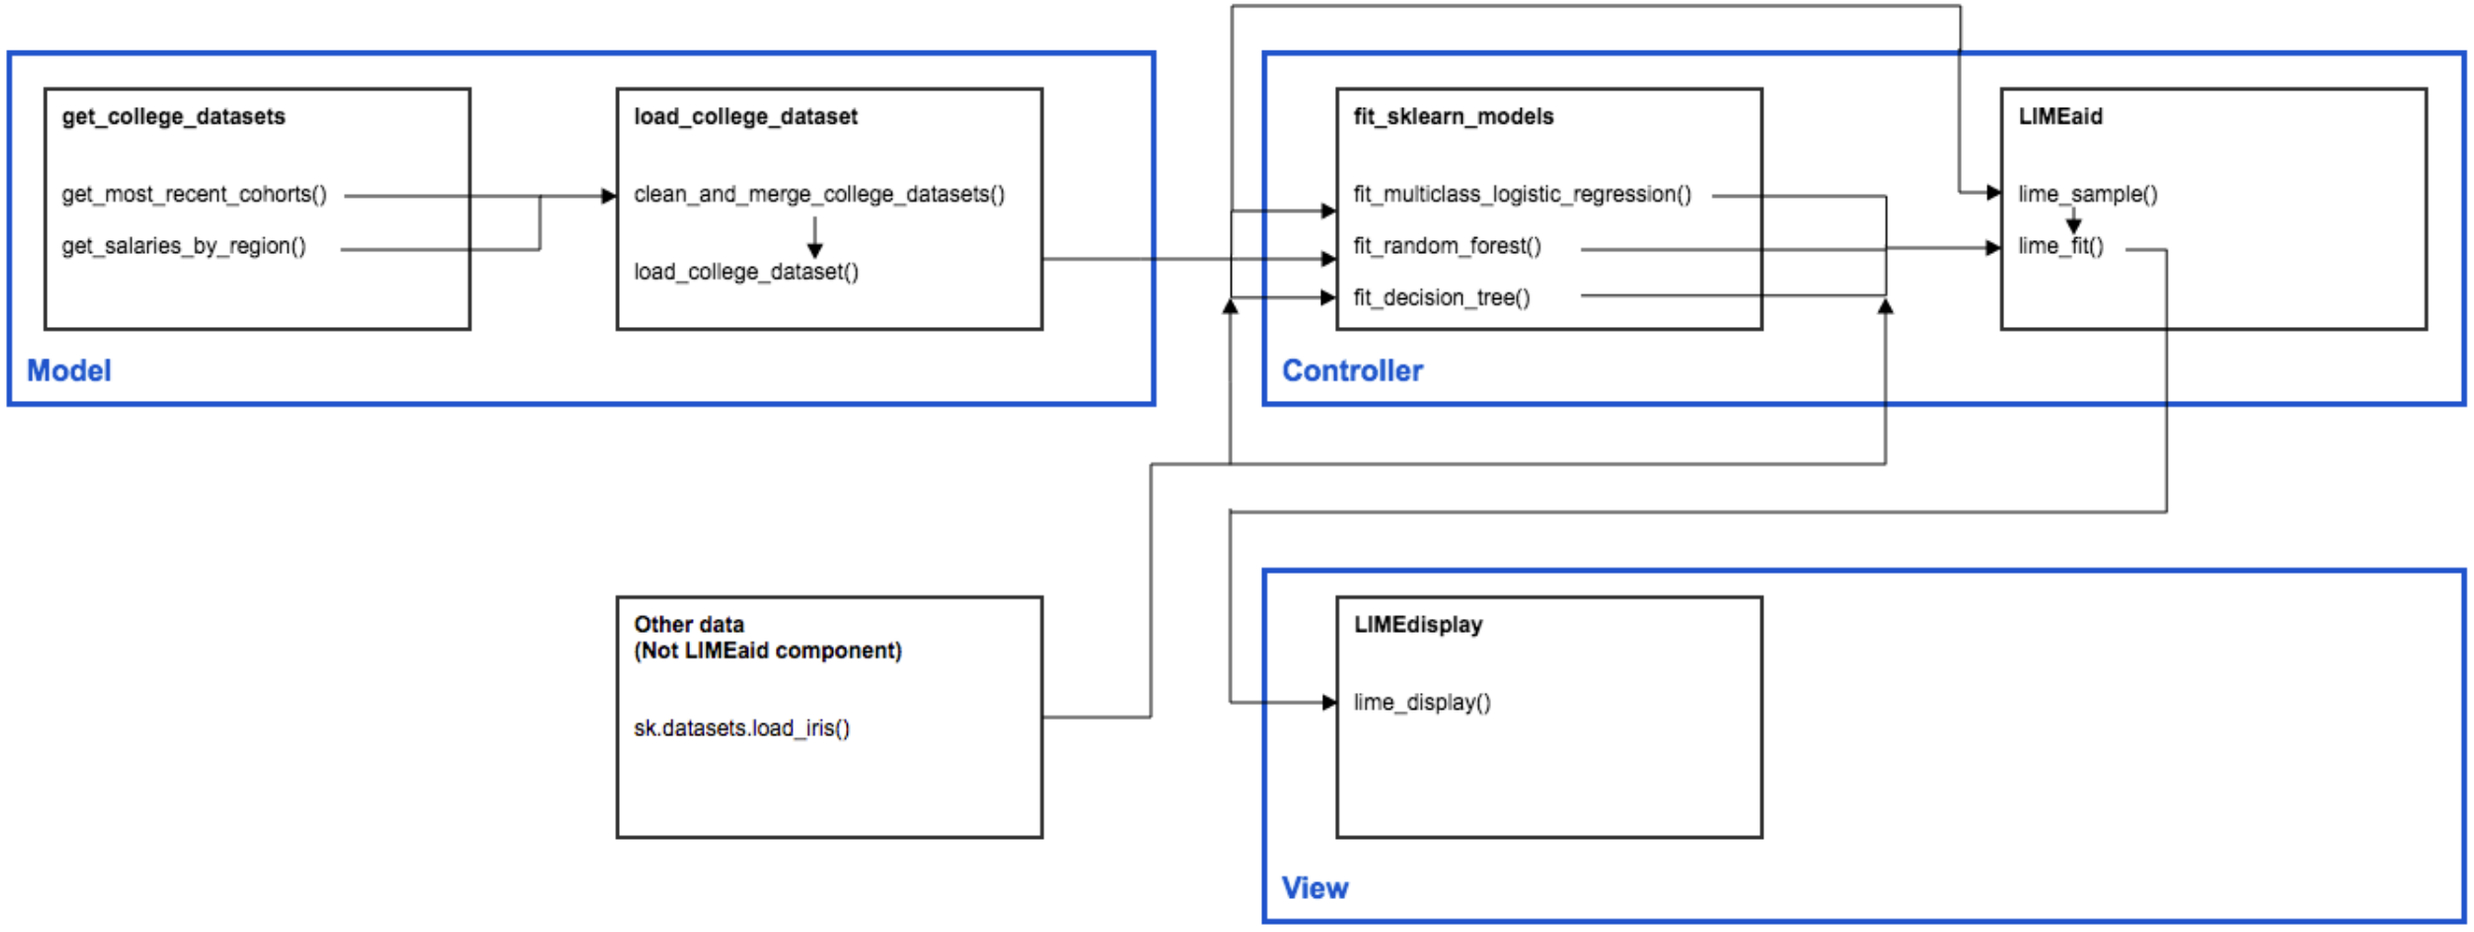
get_college_datasets
get_most_recent_cohorts()
get_salaries_by_region()
load_college_dataset
clean_and_merge_college_datasets()
load_college_dataset()
Model
fit_sklearn_models
fit_multiclass_logistic_regression()
fit_random_forest()
fit_decision_tree()
LIMEaid
lime_sample()
lime_fit()
Controller
Other data
(Not LIMEaid component)
sk.datasets.load_iris()
LIMEdisplay
lime_display()
View

Data

Sources

- College Scorecard (data.gov): Annual report of schools and attributes (SAT scores, majors offered, region, cost, public/private/for profit, etc.)
- "Where it Pays to Attend College" (Kaggle.com) obtained from (*Wall Street Journal*), based on Payscale, Inc. (College Salary Report Methodology): Article reporting salaries of graduates, salaries by major, etc.

Merge

- Significant cleaning, reformatting to match sets on college name
- String manipulation, removal of hyphens, abbreviations, region names, etc.

More

- Sklearn's provided "Iris" data

Models

Scikit-learn classifiers that predict probabilities (predict_proba implemented)

Multiclass logistic regression (sklearn.linear_model.LogisticRegression)

- 85% accuracy on College data

Random Forest (sklearn.ensemble.randomforestclassifier)

- 65% accuracy on College data

Decision tree (sklearn.tree)

- 65% accuracy on College data

Models not tuned for improved accuracy (default settings)

Demo

Education and Iris data

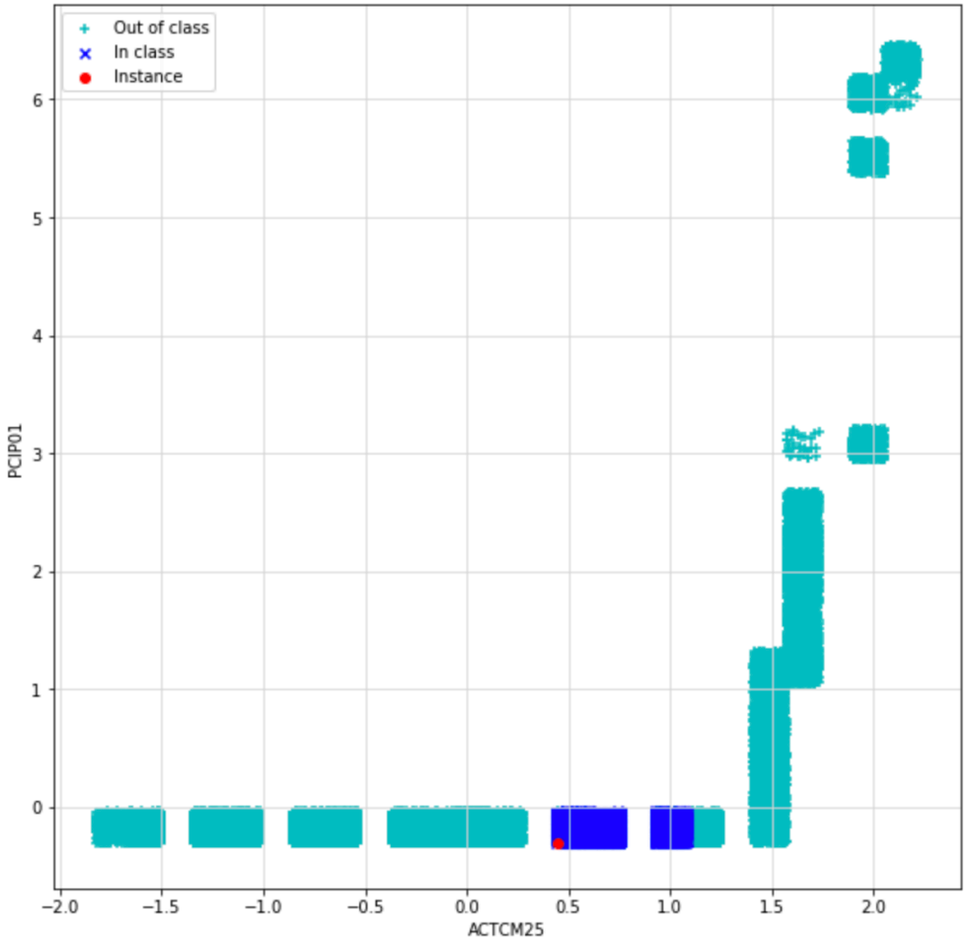

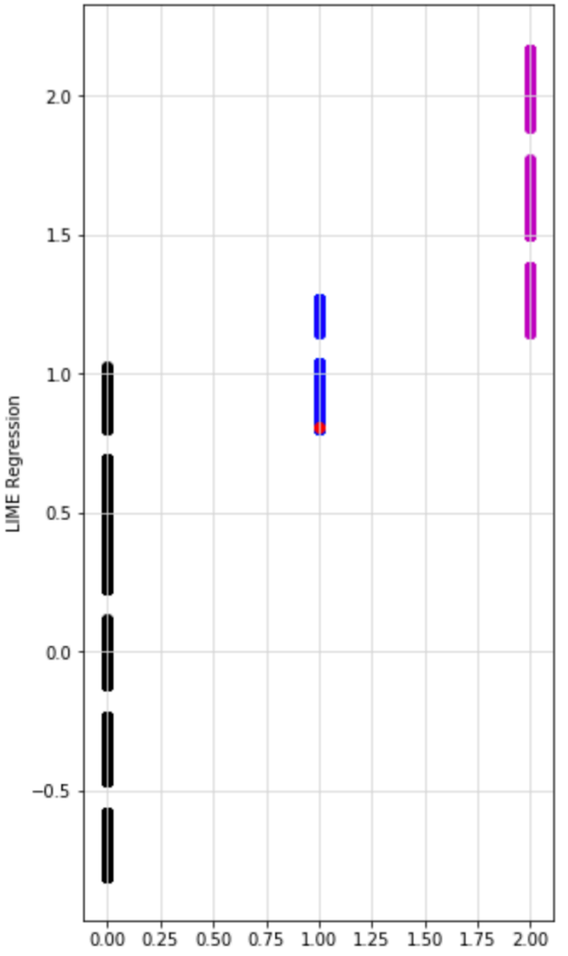

Project Structure

- LIMEaid Github Repo based on Shablona
- Used Model-View-Controller (MVC) codebase architecture

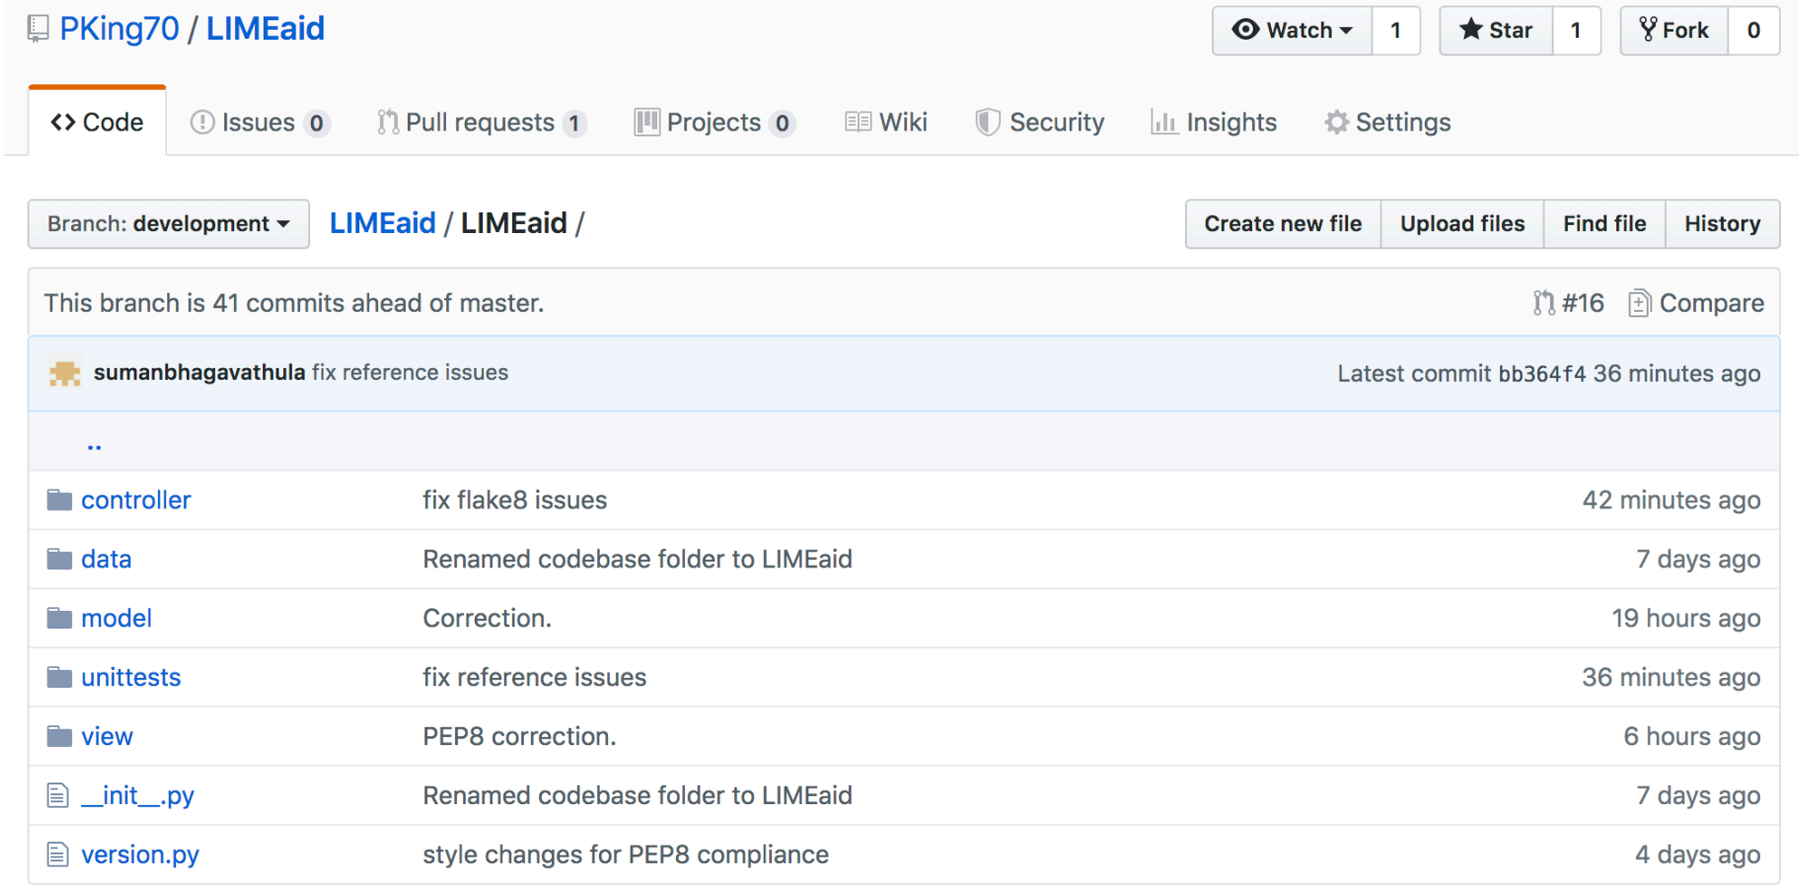

Future Work

API support for data acquisition to support dynamic features:

- College Scorecard (data.gov) currently published with new features and data dictionary yearly

Model tuning for examples:

- Currently using defaults, could improve accuracy > 65%

Modify penalty for number of coefficients

Improve method of generating random samples

More data type and model support: image data, NLP support, support for model objects beyond sklearn classifiers

Questions?